


ID1020: Union-Find

Dr. Per Brand
pbrand@kth.se

kap. 1.5



Slides adapted from Algorithms 4th Edition, Sedgewick.

Developing an algorithm

- Steps:

- Construct a model of the problem
- Find an algorithm that solves the problem.
- Is the algorithm fast enough? Memory usage ok?
- If not determine why not.
- Find a better algorithm.
- Iterate until satisfied.

- Thereafter analyze the algorithm. Determine how it scales.

Case study: Dynamic-connectivity problem

- Given N objects, the program should support the following:
 - Union**: Connect two of the objects.
 - Find**: Determine if there is path (of connections) between any two objects.
 - If there is a path then the two objects are in the same component

connect 4 and 3

connect 3 and 8

connect 6 and 5

connect 9 and 4

connect 2 and 1

are 0 and 7 connected? ✗

are 8 and 9 connected? ✓


connect 5 and 0

connect 7 and 2

connect 6 and 1

connect 1 and 0


are 0 and 7 connected? ✓



A large connectivity problem

- Is there a path between P and Q ?

Yes



Applications and the model

- Applications


- Pixels in a computer image.
- Computers in a network.
- Friends in a social network.
- Transistors on a chip.
- Elements in mathematical sets.
- Variable names in a Fortran program.

- To simplify programming, we call the objects 0 to $N - 1$.

- Use the integer as an array index.
- Abstract away the details not needed to solve/model union-find.
- $a[i] == a[j]$ for all objects in the same component


The model

- **Connectedness is an equivalence relation:**
 - Reflexive: p is connected to itself.
 - Symmetric: if p is connected to q , then q is connected to p .
 - Transitive: if p is connected to q and q is connected to r , then p is connected to r .
- An equivalence relation partition the objects into equivalence classes which in this case we call "connected components".



Find and Union operations

- **Component Identifier:** Unique id for each component
- **Find.** In what component do we find p ?
- **Connected.** Are p and q in the same component?
- **Union.** All the objects in the components that p and q become part of the same component. (In one component we replace the component identifier with that of the other component)



Union-find datatype (API)

- **Goal.** Efficient datastructure for union-find.
 - The number of objects N might be very large.
 - The number of operations M might also be very large.
 - Client can interleave find and union operations in any order.

public class **UF**

UF(int N)

*initialize union-find data structure
with N singleton objects (0 to $N - 1$)*

void union(int p , int q)

add connection between p and q

int find(int p)

component identifier for p (0 to $N - 1$)

boolean connected(int p , int q)

are p and q in the same component?

connected()

- `connected()` is implemented in one line:

```
public boolean connected(int p, int q) {  
    return find(p) == find(q);  
}
```

Example of client

- Read in the numbers objects N from stdin (standard input).
- Repeat while stdin is non-empty
 1. Read a pair of integers from stdin
 2. If not connected, create a connection (using union) and print the pair.

```
public static void main(String[] args) {  
    int N = StdIn.readInt();  
    UF uf = new UF(N);  
    while (!StdIn.isEmpty()) {  
        int p = StdIn.readInt();  
        int q = StdIn.readInt();  
        if (!uf.connected(p, q)) {  
            uf.union(p, q);  
            StdOut.println(p + " " + q);  
        }  
    }  
}
```

% more tinyUF.txt

10 ← Number of object N

4 3

3 8

6 5

9 4

2 1

8 9

5 0

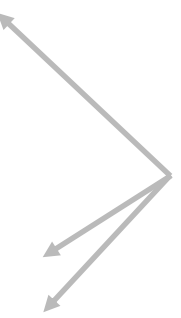
7 2

6 1

1 0

6 7

Already connected



Quick-find

Quick-find [eager method]

- Datastructure.


- Integer array `id[]` of size `N`.
- Representation: `id[p]` is the id of the component containing `p`.
`p` and `q` are connected iff they have the same id.

	0	1	2	3	4	5	6	7	8	9
id[]	0	1	1	8	8	0	0	1	8	8

0, 5 and 6 are connected

1, 2, and 7 are connected

3, 4, 8, and 9 are connected



Quick-find operations

- Find p
 - Return $id[p]$ the component identifier
- **Connected(p, q)**
 - Is $id[p] == id[q]$?

	0	1	2	3	4	5	6	7	8	9
$id[]$	0	1	1	8	8	0	0	1	8	8

$id[6] = 0; id[1] = 1$

6 and 1 are not connected

- **Union.** To unify two components (where p and q , respectively are members) change all elements $= id[p]$ to $id[q]$.


	0	1	2	3	4	5	6	7	8	9
$id[]$	1	1	1	8	8	1	1	1	8	8

↑ ↑ ↑

After union of 6 and 1


note: whole array needs to be checked

Quick-find demo



	0	1	2	3	4	5	6	7	8	9
id[]	0	1	2	3	4	5	6	7	8	9

Quick-find demo



	0	1	2	3	4	5	6	7	8	9
id[]	1	1	1	8	8	1	1	1	8	8

Quick-find: Java implementation

```
public class QuickFindUF {
```

```
    private int[] id;
```

```
    public QuickFindUF(int N)    {
```

```
        id = new int[N];
```

```
        for (int i = 0; i < N; i++) {
```

```
            id[i] = i;
```

```
        }
```

```
    }
```

```
    public boolean find(int p) { return id[p]; }
```

```
    public void union(int p, int q) {
```

```
        int pid = id[p];
```

```
        int qid = id[q];
```

```
        for (int i = 0; i < id.length; i++) {
```

```
            if (id[i] == pid)
```

```
                id[i] = qid;
```

```
        }
```

```
    }
```

```
}
```

← Initially: each object is identified by its own id (N array accesses)

← Return the id of p (1 array access)

← Change all elements with id[p] till id[q] (at most $2N + 2$ array accesses)

Quick-find is slow

- **Cost model.** The number of array accesses

algorithm	initialize	union	find	connected
quick-find	N	N	1	1

Number of array accesses


- **Union is expensive.** It will take ~~N~~ N^2 array accesses to handle of N union operations on N objekt. Quadratic!!!

Quadratic algorithms don't scale

- Rule of thumb

- 10^9 operations per second
- 10^9 memory words in main memory.
- To access entire memory takes about 1 second.

Been this way since
1950!




- Consequences for quick-find

- With 10^9 union operations on 10^9 object, quick-find will perform 10^{18} operations \Rightarrow 30+ years of computation!

- Quadratic algorithms don't scale better with technology improvements

- New computer is 10x faster.
- But, computer has 10x more memory \Rightarrow So we want to solve a problem that is 10x bigger.
- With a quadratic algorithm, \Rightarrow 10x slower



Quick Union


Quick-union [lazy method]

- Datastructure.

- Integer array `id[]` of length `N`.
- Representation: `id[i]` is parent of `i`
- Root: parent of itself
- Root of `i` is
`id[id[id[...id[i]...]]]`.

	0	1	2	3	4	5	6	7	8	9
id[]	0	1	9	4	9	6	6	7	8	9

Repeat until you reach root
(Note: not cycles)



Parent of 3 is 4

Parent of 4 is 9

9 is a root



Quick-union

	0	1	2	3	4	5	6	7	8	9
id[]	0	1	9	4	9	6	6	7	8	9

- **Find.** Determine the root of p ?
- **Connected.** Determine if p and q have the same root?
- **Union.** To connect the components that contain p and q , assign the id of p 's root to the id of q 's root.

	0	1	2	3	4	5	6	7	8	9
id[]	0	1	9	4	9	6	6	7	8	6

Only one value changes!!




Quick-union demo



	0	1	2	3	4	5	6	7	8	9
id[]	0	1	2	3	4	5	6	7	8	9

Quick-union demo



Quick-union: Java implementation

```
public class QuickUnionUF {  
    private int[] id;  
    public QuickUnionUF(int N) {  
        id = new int[N];  
        for (int i = 0; i < N; i++) {  
            id[i] = i;  
        }  
    }  
    public int find(int i) {  
        while (i != id[i]) {  
            i = id[i];  
        }  
        return i;  
    }  
    public void union(int p, int q) {  
        int i = find(p);  
        int j = find(q);  
        id[i] = j;  
    }  
}
```

← each object is initialized to be the root of a component with one member-itself (N array accesses)

← follow parent links till root is reached (number of array accesses=depth)

← make the root of p link to the root of q (number of array acceses = depth of q and p)

Quick-union is also slow

Cost model. Number of array accesses.

algorithm	initialize	union	find	connected
quick-find	N	N	1	1
quick-union	N	N [†]	N	N

← worst case

† includes cost of finding root

Quick-find defect.

- Union is too expensive. (N array accesses).
- Trees are "flat", but it costs too much to keep them *flat*.

Quick-union defect.


- Trees can become deep.
- Find/connected/union too expensive (may take N array accesses).

Quick-union improvements

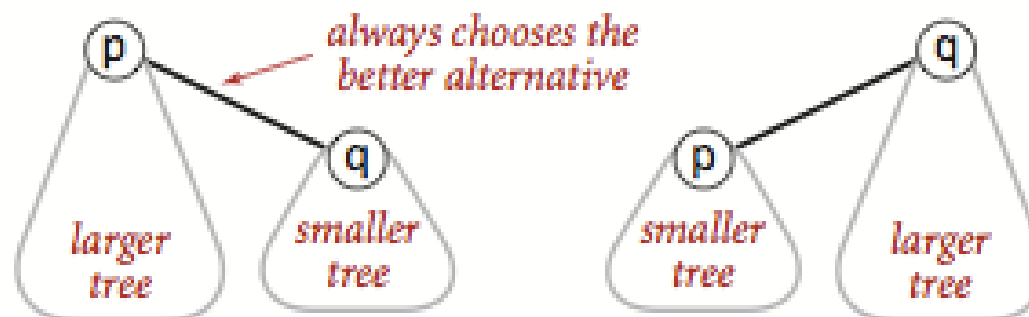
Improvement 1: Weights

- Weighted quick-union.
 - Change quick-union to avoid deep trees.
 - Store size of tree (number of objects).
 - Balance tree with linking the root of the smaller tree to the larger tree.

quick-union



weighted




Weighted quick-union

Viktad quick-union demo



	0	1	2	3	4	5	6	7	8	9
id[]	0	1	2	3	4	5	6	7	8	9


Viktad quick-union demo



	0	1	2	3	4	5	6	7	8	9
id[]	6	2	6	4	6	6	6	2	4	4

Quick-union and weighted quick-union examples

quick-union



average distance to root: 5.11

weighted



average distance to root: 1.52

Quick-union and weighted quick-union (100 sites, 88 union() operations)

Weighted quick-union: Java implementation

- **Datastructure.** Like quick-union, but we use an extra array `sz[i]` to keep track of number of object in tree
- **Find/connected.** The same as for quick-union.
- **Union.** Changes:
 - Link the root of the smaller tree to the larger.
 - Update `sz[]` array.

```
int i = find(p);
int j = find(q);
if (i == j) return;
if (sz[i] < sz[j]) {
    id[i] = j; sz[j] += sz[i];
} else {
    id[j] = i; sz[i] += sz[j];
}
```


Weighted quick-union analysis

- Running time.

- Find: time proportional to the depth of p .
- Union: takes constant time, given the roots.

- Theorem. Depth of object is at most $\lg N$.

\lg = base-2 logarithm




$$N = 11$$

$$\text{djupet}(x) = 3 \leq \lg N$$

Weighted quick-union analysis

- **Theorem.** Depth of object is at most $\lg N$.
- **Proof.** What can increase the depth of node x ?
 - Increase by 1 when the tree T_1 containing x is merged with another tree T_2 .
 - Size of tree containing x is at least doubled as $|T_2| \geq |T_1|$.
 - Size of tree containing x can be doubled at most $\lg N$ times. Why?



Weighted quick-union analysis

- Running time.

- Find: time proportional to the depth of p i.e at most $\lg N$.
- Union: takes constant time, given the roots. $\sim \lg N$


algorithm	initialize	union	find	connected
quick-find	N	N	1	1
quick-union	N	N^\dagger	N	N
weighted QU	N	$\lg N^\dagger$	$\lg N$	$\lg N$

\dagger including cost of finding root


Can the algorithm be further improved. Yes, a bit.

Improvement 2: path compression


- Quick union with *path compression*. Flatten the tree when moving to find root.




Improvement 2: path compression



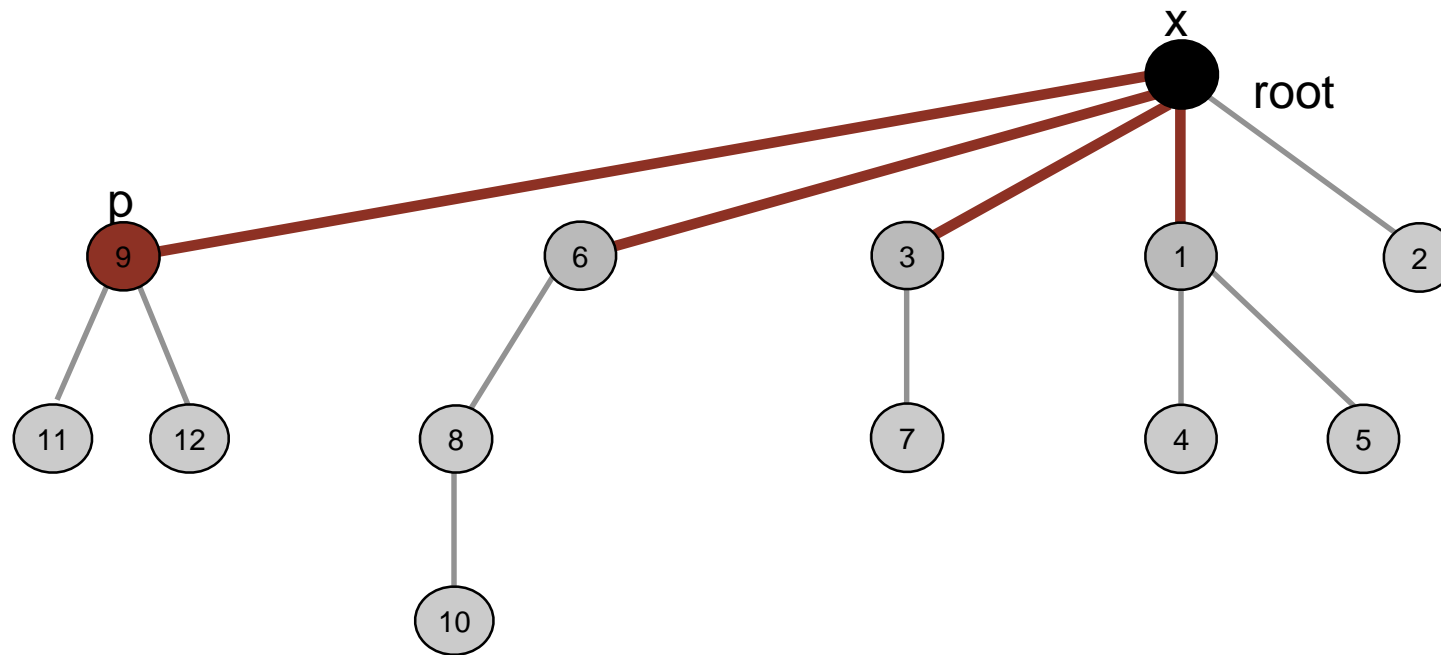
Improvement 2: path compression



Improvement 2: path compression



Improvement 2: path compression



Path compression: Java implementation

- **Two step implementation:** add an extra loop to `find()` to assign new root `id[]` to every examined root.

```
public int find(int i) {  
    while (i != id[i]) {  
        id[i] = id[id[i]];  
        i = id[i];  
    }  
    return i;  
}
```

← One extra line of code

Note. Path is compressed as the cost of increased constant overhead in certain operations.

Path compression even better?

No linear algorithm for M union-find operationer on object!!

- In theory WQUPC (weighted quick-union with path compression) is not linear
- But in practice it is WQUPC linjär.

Conclusion

- **Weighted quick union** (with or without path compression) enables problems to be solved that otherwise could not be

algorithm	worst-case tid
quick-find	$M N$
quick-union	$M N$
viktad QU	$N + M \log N$
QU + stig komprimering	$N + M \log N$
viktad QU + stig komprimering	Almost but not quite $N + M$

Runtime for M union-find operationer on N objects


Ex. [10^9 unions and finds with 10^9 objects]

- WQUPC diminishes running time from 30 years to 6 seconds.
- Faster computers don't help – good algorithms do

Union-Find applications


Union-find applications

- Perkolations.
- Spel (Go, Hex).
- ✓ **Dynamisk konnektivitet.**
- Least common ancestor.
- Ekvivalens av ändliga tillståndsautomater.
- Hinley-Milner polymorphic type inference.
- Kruskal's spanning tree algorithm.
- Matlab's `bwlabel()` function i image processing.




Percolation

- An abstract model of many physical systems
 - $N \times N$ grid of *sites*.
 - A site is open with probability p (and blocked with probability $1 - p$).
 - The system will **percolate** iff top and bottom are connected via open sites.



$N = 8$




Percolation

- An abstract model of many physical systems
 - $N \times N$ grid of *sites*.
 - A site is open with probability p
(and blocked with probability $1 - p$).
 - The system will **percolate** iff top and bottom are connected via a chain of open sites


model	system	vacant site	occupied site	percolates
electricity	material	conductor	insulated	conducts
fluid flow	material	empty	blocked	porous
social interaction	population	person	empty	communicates

Percolation


- Depending on grid size N and *site vacancy* probability p .






p low (0.4)
does not percolate



p medium (0.6)
percolates?



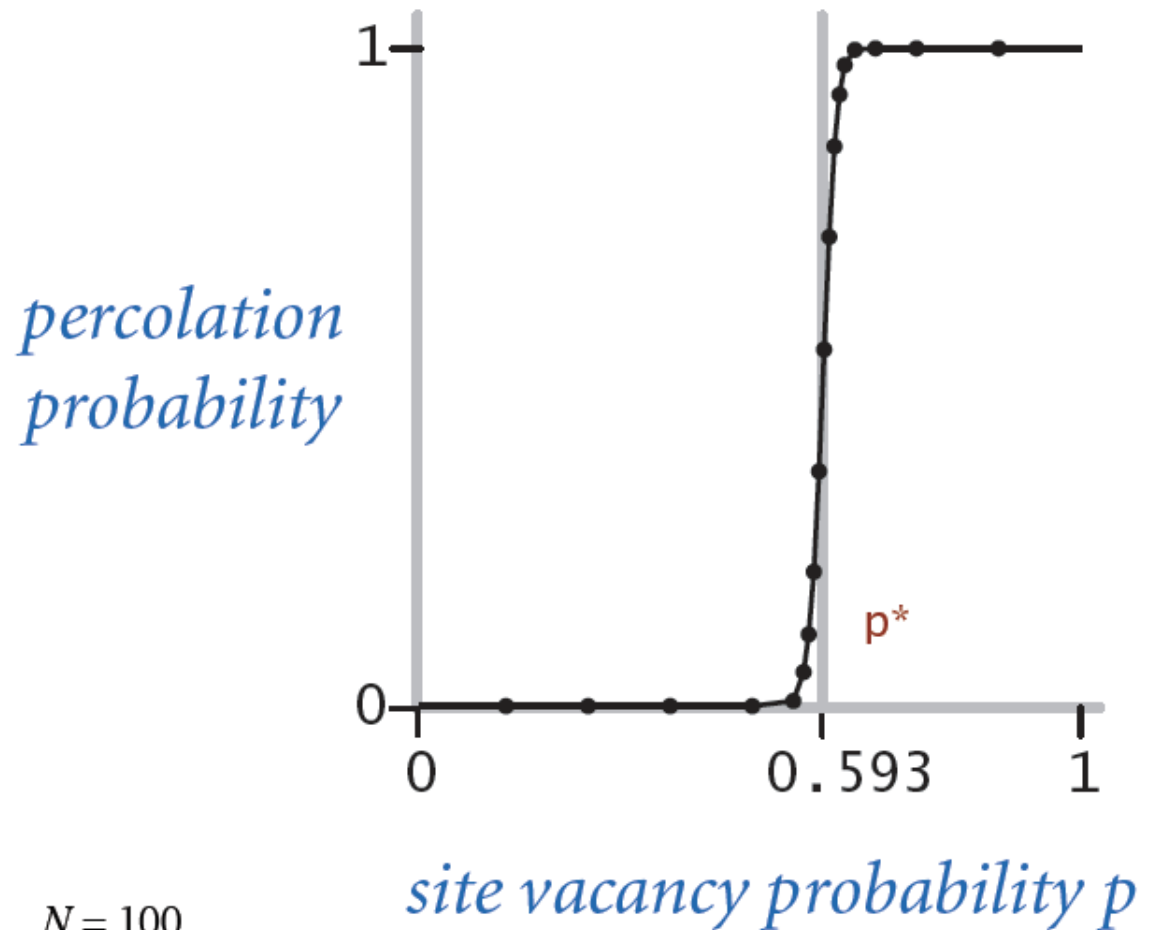
p high (0.8)
percolates



Percolation phase transition


- When N is large, theory guarantees a sharp threshold p^* .
 - $p > p^*$: percolates with high probability.
 - $p < p^*$: does not percolate with high probability.

- What is the value of p^*




Monte Carlo simulating

- Initialize all sites in a $N \times N$ grid to be blocked
- Make randomly chosen sites open until there is a path between top and bottom.
- Determine *vacancy percentage*, determine p^* .



Model as a connectivity problem


- How to check if a $N \times N$ system percolates?
Modellera as a **dynamic connectivity** problem and use **union-find**.






Model as a connectivity problem

- ✓ Create an object for each site and name them from 0 to $N^2 - 1$.

$N = 5$




 Open site
 Blocked site



Model as a connectivity problem


- How to check if a $N \times N$ system percolates?
 - ✓ Create an object for each site and name them from 0 to $N^2 - 1$.
 - ✓ Sites are in the same component iff they are connected by open sites.

$N = 5$



 Open site


 Blocked site



Model as a connectivity problem


- How to check if a $N \times N$ system percolates?
 - ✓ Create an object for each site and name them from 0 to $N^2 - 1$.
 - ✓ Sites are in the same component iff they are connected by open sites.
 - ✓ System percolates iff any site on the bottom row is connected to some site on the top row

brute-force algorithm: N^2 calls to `connected()`



 Open site



 Blocked site



Improvement



- **Trick.** Add 2 virtual sites that have connections to all sites in top and bottom rows, respectively.
- ✓ System percolates iff the virtual top site is connected to the virtual bottom site

mer effektivt algoritm: bara 1 anrop till connected()



Blocked to open site



- How do we change a site from blocked to open?



Blocked to open site

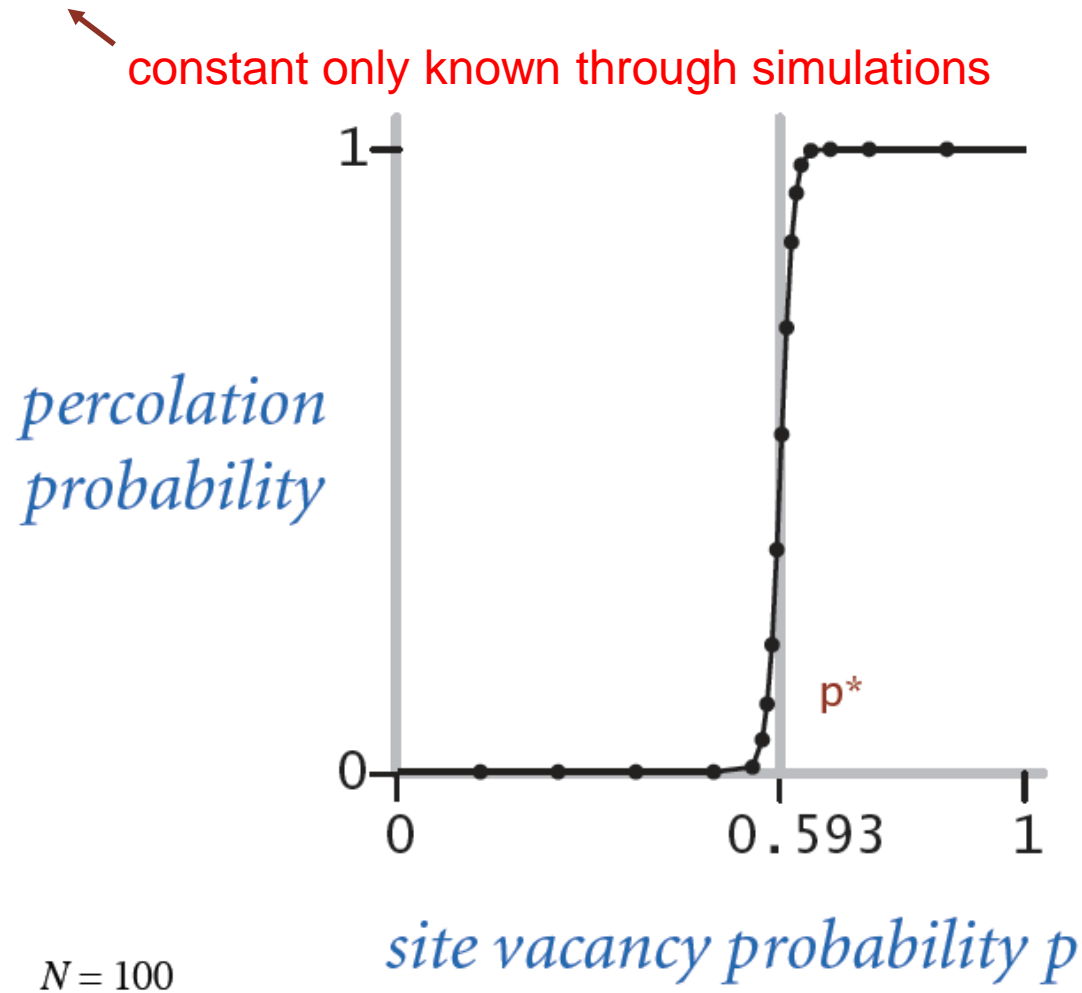
- How do we change a site from blocked to open?
- Connect the site to its 4 neighbors .

4 calls to union()



Percolation Threshold

- What is the percolation threshold p^* ?
- Approx. 0.592746 for *large square lattices*.



- A fast algorithm enables an accurate answer to a scientific question.

Summary

- Union-Find is an interesting class of algorithm with practical applications
 - Quick-Find, Quick-Union, Weighted Quick-Union, Weighted Quick-Union with Path Compression.
- Using union find as a case study, we have covered the stages of algorithm development:
 - Build a model of the problem
 - Find an algorithm to solve the problem
 - Evaluate. Is it fast enough? Will it fit in memory?
 - If not, determine why?.
 - Find a better algorithm.
 - Iterate until satisfied.